

Computer Architecture

Performance

- X is n times faster than Y

$$\frac{Perf(X)}{Perf(Y)} = \frac{Execution\ time(Y)}{Execution\ time(X)} = n$$

$$CPU\ time = \frac{\#cycles/program}{frequency\ (Hz)}$$

$$F_c = \frac{\#type\ instr}{prog\ instr\ count}$$

$$avg\ cycle/instr\ CPI = \sum_{i=1}^n CPI_i \cdot F_i$$

$$CPU\ execution\ time = \frac{(instr\ count)}{frequency\ (Hz)} \cdot (avg\ CPI)$$

$$MIPS = \frac{instr\ count}{execution\ time \times 10^6}$$

- benchmark → standardized tests

CPU Power

$$P = VI$$

$$= (capacitance\ load)(voltage)^2(freq)$$

- faster CPU = more heat

→ reduce heat by (cooling, ↓ voltage, ↓ leaked current, ↓ clockspeed) → ↓ performance

Instructions

→ R type - deals w values in registers

→ I type - memory access (load)

→ S type - memory access (store)

→ SB type - branch instrs (2¹² range)

→ U type - upper immediates

→ UJ type - jump instrs

→ See "Commands" for formats

Registers

x0 - constant 0

x1 - return address

x2 - stack pointer

x3 - global pointer

x4 - thread pointer

x5-x7, x28-x31 - temporary register

x8 - frame pointer

x9, x18-x27 - saved register

x10-x17 - function argument/result

Addressing

low Addr

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31

Byte 0 1 2 3 4 5 6 7 Little endian

Byte 7 6 5 4 3 2 1 0 Big endian

Half → 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15

32 → 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31

- Arrays

→ A[i] → base + offset

$$offset = (type\ size)(index)$$

→ A[4] 8 bit → (8)(4) = 32

$$32 + base\ addr = A[4]$$

- can shift to adjust for type size

Machine Language → RISC-V

func7 rs2 rs1 fun3 rd opcode R

imm[11:0] rs1 rs2 rd opcode I

[11:5] rs2 rs1 rs3 [4:0] opcode S

[11:0] rs2 rs1 rs3 [11:0] opcode B

imm[31:12] rd opcode U

imm[31:12] rd opcode J

Control Flow

- Lui tip
→ load 0xDEADBEEF to x10
lui x10, 0xDEADBEEF
→ 1/2 bit value B=1011 starts w 1
addi x10, x10, 0x00000001
→ b/c lui writes to 31:12
sign extends 63:32 &
clears 11:0

Control Flow

→ decision making instrs

Ex Loop:

<Procedure>
bne x1, x2, Exit
<Procedure if not Exit>
→ bge x0, x0, Loop
Exit: ...
guarantees no error thrown
(may cause ∞ loop though)

- branch instr are SB type

→ range 2¹² bits

- location addr's must be even #s if given odd

① subtract 1

② raise exception/error

- for moving bwn instrs

→ jal jump & link, go to procedure

2⁶ bit reach

→ jalr jump & link register, return to previous branch

- when jumping procedures

① Store params in stack

(increment stack pointer)

② transfer to procedure + preform ops

③ update values in stack for caller to use

④ return to place of call (jalr)

Ch 3 3 Ch 4

Multiplication

→ m bits x n bits = m+n bit product

→ Algo 1

① 0's (same as multiplier) multiplicand

Each stage

② if LSB of multiplier = 1

product + multiplicand

if 0 → do nothing

③ shift multiplier right

④ shift multiplicand left

Cont to repeat until multiplier gone

Ex 1001₂ × 1001₂

0 1001 0001001 00000000

1 0100 0001001 00001001

2 0010 00100100 00001001

3 0001 01001000 00001001

4 0000 10010000 01010001

ans

- max value → 64 bit × 64 bit

- 1 cycle/step (3 steps) (64 bits) × 200 clocks

→ Optimized Algo

① add multiplier to product

Each stage

② if LSB of product = 1,

add multiplier to MSB

if 0 → do nothing

③ shift product right

Cont to repeat until multiplier gone

Ex 1001₂ × 1001₂

0 1001 0001001 00000000

1 1001 0001001 00001001

2 0100 00100100 00001001

3 0001 01001000 00001001

4 0000 10010000 01010001

ans

Booth's Algo (Log(4))

→ signed multiplication

→ mule both #'s +,

→ 2's comp any neg

* sign extend

① Partition

② if 01 → add multiplicand

if 10 → sub multipli

③ shift right

Ex 2 × -3

multiplicand 2 = 0010

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

2's comp = 1110

significand base
accuracy/fraction (1.x) × 2^{exponent}
In N bits, range
→ unsigned 0 to 2^N
→ 2's comp -2^{N-1} to 2^{N-1}
→ signed ±2^{N-1}

Floating pt larger range

→ half (16 bit) 2⁻¹⁴ to (2⁻²)2¹⁵

→ single (32) 2⁻¹²⁶ to 2¹²⁷

→ double (64) 2⁻¹⁰²² to 2¹⁰²⁴

sign exponent fraction

16 = 1 5 10 half

32 = 1 8 23 single

64 = 1 11 52 double

→ x = (-1)^{sign} (1 + frac) (2^{exp - bias})

bias → 15, 127, 1023, = 2^(#exponents-1)

-denormal → hidden bit = 0, 0.x

-∞ → exp = 11...11 frac = 00...00

-NAN → exp = 11...11 frac ≠ 00...00

Bin → dec

1 10000011 111000... →

(-1)¹ (1 + 2⁻¹ + 2⁻² + 2⁻³), 2⁽¹²⁷⁻¹²⁶⁾

= -30

Normalize 1010 → 1.010 × 2³

Dec → bin

sign → 1

Exp → 3 + bias = 130₁₀ = 10000010₂

Frac → part after decimal 010

→ 1 10000010 01000000...

Smallest exp 00...001 = 1

largest exp 11...110 value depend on type

Dec → bin -0.75

sign → 1

Exp → -1 + bias = 126₁₀ = 01111110₂

Frac → 100...00

→ 1 01111110 1000...00

OR

0.75 × 2 = 1.5

0.5 × 2 = 1

0.25 × 2 = 0.5

0.125 × 2 = 0.25

0.0625 × 2 = 0.125

0.03125 × 2 = 0.0625

0.015625 × 2 = 0.03125

0.0078125 × 2 = 0.015625

0.00390625 × 2 = 0.0078125

0.001953125 × 2 = 0.00390625

0.0009765625 × 2 = 0.001953125

0.00048828125 × 2 = 0.0009765625

0.000244140625 × 2 = 0.00048828125

0.0001220703125 × 2 = 0.000244140625

0.00006103515625 × 2 = 0.0001220703125

0.000030517578125 × 2 = 0.00006103515625

0.0000152587890625 × 2 = 0.000030517578125

0.00000762939453125 × 2 = 0.0000152587890625

0.000003814697265625 × 2 = 0.00000762939453125

0.0000019073486328125 × 2 = 0.000003814697265625

0.00000095367431640625 × 2 = 0.0000019073486328125

0.000000476837158203125 × 2 = 0.00000095367431640625

0.0000002384185791015625 × 2 = 0.000000476837158203125

0.00000011920928955078125 × 2 = 0.0000002384185791015625

0.000000059604644775390625 × 2 = 0.00000011920928955078125

0.0000000298023223876953125 × 2 = 0.000000059604644775390625


0.00000001490116119384765625 × 2 = 0.0000000298023223876953125

Processors

- 3 stages
 - ① PC supplies instr addr & fetches instr from memory (& update PC)
 - ② decode instr & read regs
 - ③ execute instrs

- * All instr except use ALU
- 2 functional units
 - ① Combinational → operate on data
 - ② Sequential → contain state

- Clocking methodology




- cycle length = longest logic gate delay
- instr fetch
 - ① read PC addr
 - ② pull from instr & execute
 - ③ $PC + 4$
- R format instr
 - ① read 2 reg ops
 - ② do arithmetic/logic ops
 - ③ (load) read mem & update reg
 - ④ (store) write reg value to mem
- branch instr
 - ① read reg ops
 - ② compare ops
 - ③ calculate target addr
 - ④ sign extend / shift / $PC + 4$
- control unit → regulates data transfer
- decodes instrs (w/ instr type & func)
- ALU input → $A' + B' + op$

0	0	0	0	and
0	0	0	1	or
0	0	1	0	add
0	1	1	0	sub
1	1	0	0	set less than
1	1	0	1	nor

Pipelining

- 5 stages
 - ① Fetch (instr) & update PC
 - ② dec - reg read & instr decode
 - ③ exec - calculate mem addr
 - ④ mem - read data from mem
 - ⑤ wb - write back data to reg

- increases efficiency
- executes multiple instrs @ same time



Hazards

- situations preventing starting the next instructions
- Control Hazard
 - deciding ctrl action depends on a prev instruction
 - SLTN: branch prediction
 - execute a possible branch & only stall if guess is wrong
 - static → uses typical behavior
 - dynamic → table of previous decisions
- Structure Hazards
 - required resource is busy, really only an issue w/ instr & data
 - SLTN? have separate instr & data memory paths
- Data Hazards
 - need data update from a prev instr
 - SLTN: forwarding (pull value after exec phase / don't wait for write back)

add	IF	Dec	Exec	Mem	WB
sub	IF	Dec	Exec	Mem	WB
 - doesn't work for ld (bc gets value @ end of mem)
 - so insert 1 stall/bubble/NOP
- Can also reorder code to minimize hazards